

Deep Learning Image Classification Project Report

Introduction

This report details the implementation and comparison of three approaches for image classification using the CIFAR-10 dataset: a custom CNN architecture and two transfer learning approaches using VGG16 (with and without fine-tuning). The project aimed to evaluate different methodologies and their effectiveness in classifying images across 10 categories.

Methodology & Architecture Decisions

1. Custom CNN Architecture

We implemented a VGG-style architecture with the following key design decisions:

Progressive Depth:

- Started with 64 filters and doubled at each block (64→128→256)

Regularization Techniques:

- L2 regularization (0.01) to prevent overfitting
- Dropout rates (0.4→ 0.5) for deeper layers
- BatchNormalization after each convolutional layer

Optimization Strategy:

- Adam optimizer with initial learning rate of 0.001
- Learning rate scheduling for better convergence

IMAGE 1: Architecture Diagram showing the custom CNN structure

```
model = Sequential([
    Conv2D(64, (3,3), padding='same', activation='relu', input_shape=(32,32,3), kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    Conv2D(64, (3,3), padding='same', activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    MaxPooling2D(),
    Dropout(0.4),

    Conv2D(128, (3,3), padding='same', activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    Conv2D(128, (3,3), padding='same', activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    MaxPooling2D(),
    Dropout(0.5),

    Conv2D(256, (3,3), padding='same', activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    Conv2D(256, (3,3), padding='same', activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    MaxPooling2D(),
    Dropout(0.5),

    Flatten(),
    Dense(512, activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    Dropout(0.3),
    Dense(10, activation='softmax')
])
```

2. Transfer Learning with VGG16

We chose VGG16 for transfer learning because:

- Proven architecture for image classification
- Well-documented performance on similar tasks
- Relatively simple architecture compared to more recent models

Two approaches were implemented:

Basic Transfer Learning:

- Frozen VGG16 layers
- Custom classification head
- Global Average Pooling to reduce parameters

Fine-tuned Transfer Learning:

- Unfroze last 4 layers of VGG16
- Lower learning rate (0.0001)
- Same custom classification head

Results Analysis

Performance Metrics Comparison




IMAGE 2: Bar chart comparing metrics across models

Key Model	Accuracy	F1-Score	Precision	Recall
Custom CNN	88%	0.88	0.88	0.88
VGG16 Transfer	63%	0.638	0.637	0.640
VGG16 Fine-tuned	75%	0.754	0.755	0.754

Findings

Custom CNN Performance:







- Best overall performance (80% accuracy)
- Consistent metrics across all evaluation criteria
- Suggests better adaptation to CIFAR-10's specific characteristics
- Simple image augmentations were performed but no significant improvement was observed.

Basic Transfer Learning Limitations:

- Unexpectedly lower performance (63% accuracy)
- Possible reasons:
 - VGG16's original training on higher resolution images
 - Architecture complexity vs. CIFAR-10's simple nature
 - Feature mismatch between ImageNet and CIFAR-10

Fine-tuning Improvements:

- Significant improvement over basic transfer learning
- 12% accuracy increase through fine-tuning
- Shows importance of adapting pre-trained weights



Discussion

1. Why Custom CNN Performed Better

- **Tailored Architecture:** Designed specifically for 32x32 images
- **Appropriate Complexity:** Balanced depth and width for CIFAR-10
- **Effective Regularization:** Combined dropout, L2, and BatchNorm

2. Transfer Learning Challenges

- **Resolution Mismatch:** VGG16 optimized for 224x224 images
- **Feature Abstraction:** ImageNet features may be too complex
- **Model Capacity:** Possibly overparameterized for CIFAR-10

3. Lessons Learned

Architecture Fit:

- Simpler isn't always worse
- Match model complexity to data complexity

Transfer Learning Considerations:

- Pre-trained models need careful adaptation
- Fine-tuning is crucial for performance
- Domain similarity matters

Optimization Insights:

- Learning rate scheduling improved convergence
- BatchNormalization crucial for training stability
- Progressive dropout rates effectively managed overfitting

Conclusion

The custom CNN emerged as the most effective solution for this specific task, challenging the common assumption that transfer learning always provides better results. This highlights the importance of considering dataset characteristics and model complexity when choosing an architecture. The project demonstrates that while transfer learning is powerful, a well-designed custom architecture can sometimes be more appropriate for specific use cases.

Future Work

- 1. Experiment with modern architectures (ResNet, EfficientNet)
- 2. Implement data augmentation strategies
- 3. Explore ensemble methods combining different approaches
- 4. Investigate the impact of image resolution on transfer learning

IMAGE 4: Confusion matrices
Show confusion matrices for all three models to visualize class-wise performance

CNN Model Metrics

```
Matriz de Confusión:
[[892 10 20 5 9 2 7 3 34 18]
 [ 3 947 1 1 0 0 2 1 6 39]
 [ 35 0 813 27 39 28 40 11 5 2]
 [ 13 5 34 753 23 97 41 15 11 8]
 [ 4 0 24 20 898 8 19 23 4 0]
 [ 4 2 18 67 25 846 11 22 1 4]
 [ 1 1 16 19 13 11 934 2 2 1]
 [ 6 0 6 11 28 21 5 919 1 3]
 [ 24 8 2 3 2 1 1 0 945 14]
 [ 9 32 1 0 0 1 3 0 12 942]]

Métricas de Evaluación:
F1-Score: 0.8882
Precisión: 0.8883
Recall: 0.8889
```

VGG16 Model Metrics

```
Matriz de Confusión:
[[728 27 44 20 13 12 11 20 87 38]
 [ 27 721 11 40 9 17 19 14 31 111]
 [ 61 16 537 72 106 52 96 35 10 15]
 [ 19 36 65 465 60 155 99 31 22 48]
 [ 18 13 78 55 580 40 100 77 27 12]
 [ 11 15 52 189 60 535 45 56 9 28]
 [ 6 17 60 70 57 44 718 3 11 14]
 [ 22 13 35 44 64 64 11 705 4 38]
 [ 71 54 19 17 11 6 8 8 763 43]
 [ 38 121 12 34 13 16 16 32 48 670]]

Métricas de Evaluación:
F1-Score: 0.6412
Precisión: 0.6411
Recall: 0.6422
```

Fine tune Model Metrics

```
Matriz de Confusión:
[[814 16 40 17 13 3 7 7 64 19]
 [ 16 868 1 8 3 6 9 2 18 69]
 [ 30 7 729 55 58 27 61 17 7 9]
 [ 14 16 51 591 45 137 76 27 16 27]
 [ 17 9 90 56 676 25 57 52 10 8]
 [ 11 8 27 214 39 612 34 39 0 16]
 [ 4 10 42 54 32 24 817 0 10 7]
 [ 20 5 27 37 46 46 11 792 0 16]
 [ 43 40 10 12 12 1 5 3 854 20]
 [ 31 103 6 18 5 7 8 9 20 793]]

Métricas de Evaluación:
F1-Score: 0.7542
Precisión: 0.7551
Recall: 0.7546
```